

# GPU Programming

(in Cuda)

---

Julian Gutierrez

NUCAR

Session 1

# Outline

- Introduction to the Course
- Syllabus
- Introduction to Parallel Programming


# Brief Introduction

# Who am I

- Julian Gutierrez
  - [jgutierrez@ece.neu.edu](mailto:jgutierrez@ece.neu.edu)
- Invited Lecturers
  - Leiming Yu (PhD)
  - Shi Dong (PhD)
  - Charu Kalra (PhD)

# NUCAR


- Northeastern University Computer Architecture Research laboratory
- Under the direction of Dr. David Kaeli
- We do research in many computer engineering areas:
  - Embedded Systems
  - Architecture Simulation
  - GPU Computing
  - Machine Learning
  - Many cores
  - Reliability
  - Security
  - And others



NUCAR


# Who are you

## C/C++ Programming Experience



- None
- Basic
- Medium
- Advance


## CUDA Programming Experience




- None
- Basic
- Medium
- Advance

# Who are you

## Main Program of Study



## Level of Study



# What the class is about

- The main idea
  - Free class where you can learn and interact with GPUs
  - Learn how to write and design algorithms to develop efficient and high performance programs on GPUs
- We will offer a certification for those students who participated in the course as well as a certification for the best performing program for the final project.



# Schedule

- The course consists of ~12 sessions, 2 sessions per week.

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
September						
					1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30
October						
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

# Schedule

- 6:00 pm to 7:30 pm on Mondays at 221 Hayden Hall
- 4:40 pm to 6:10 pm on Wednesdays at 011 Kariotis Hall

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
September						
					1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30
October						
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

# Class structure

Class	Topics
1	Introduction / Discovery Cluster / Basics
2	Cuda Execution Model / Architecture and NVVP and NVPROF
3	<b>Lab 1</b>
4	Cuda Memory Model
5	Optimizing Memory Performance and Synchronization
6	<b>Lab 2</b>
7	Advanced Memory Topics: Pinned, Unified. Concurrency and Dynamic Parallelism.
8	<b>Lab 3</b>
9	Image Processing
10	<b>Lab 4</b>
11	Applications (Parallel Reduction, Scan, Histogram, Convolution)
12	Additional Topics (OpenACC, Modern GPU architecture, Pascal, CUDA 8, CUDNN) and Final Project Presentation

# Class structure

Class	Topics
1	Introduction / Discovery Cluster / Basics
2	Cuda Execution Model / Architecture and NVVP and NVPROF

**NOT DEFINITIVE!**

9	Image Processing
10	<b>Lab 4</b>
11	Applications (Parallel Reduction, Scan, Histogram, Convolution)
12	Additional Topics (OpenACC, Modern GPU architecture, Pascal, CUDA 8, CUDNN) and Final Project Presentation

# How the class is going to work

- Lets have fun, learning experience for all of us
- Bring computers to class. Some days we will be doing hands on labs
- To receive final certificate you are required to attend all labs

# Final Project

- There will be one small final project. The project will consist of a small competition at the end of the course where all students will have to improve the performance of an already existing algorithm provided by the professor.
  - Project Start: Image Processing Session
  - Project Deadline: Last Session
  - Field: Image processing
- The fastest (on average) to provide a correct output will receive a certificate of champion.

# Piazza

- All class resources will be found in piazza
- Please register in the group to be able to access the resources

[www.piazza.com/northeastern/fall2017/nugpu101](http://www.piazza.com/northeastern/fall2017/nugpu101)

# Discovery Cluster

- The server that will be used for the class will be the discovery cluster, having Dr. David Kaeli as the Sponsor
- You will have access to such server to work on the project



Take out your  
computers / phones /  
tablets / calculators /  
anything with access to  
the internet

# How to setup an account for the Discovery Cluster

- In case you don't have an account on discovery cluster:
  - Discovery cluster main page:  
<http://www.northeastern.edu/rc/>
  - Go to:  
[https://www.northeastern.edu/rc/?page\\_id=20](https://www.northeastern.edu/rc/?page_id=20)
    - Fill form.
    - If you use the pdf version. Open with Adobe Reader (Chrome/web browsers don't store modifications to the pdf files).
    - Submit to [researchcomputing@neu.edu](mailto:researchcomputing@neu.edu)

# How to setup an account for the Discovery Cluster

- Users Full Name
  - Your name
- Organization and Department
  - Your main studies department name (E.g. Computer Engineering)
- Email
  - Use husky email
- Designation
  - Grad Student / Student
- Northeastern University Username
  - The username of your husky account
- Sponsor Name
  - David Kaeli
- Sponsor Email
  - [kaeli@ece.neu.edu](mailto:kaeli@ece.neu.edu)
- Type of Research
  - NUCAR GPU Programming Class
- Details of intended use
  - Will be using GPU nodes to learn how to program in CUDA.
- **Accept User Policy**
- **Save pdf/document and make sure it stored everything**
- **Submit to [researchcomputing@neu.edu](mailto:researchcomputing@neu.edu)**


# How to connect to the Discovery Cluster

- Need an SSH connection.
  - I recommend using Unix Terminal.
  - If using Windows, you could use PuTTY or similar.
- We will cover this in our first lab.

# Introduction to Parallel Programming


# The move to multi-core computing

- The CPU industry has elected to jump off the cycle-time scaling bandwagon
  - Power/thermal constraints have become a limiting factor
  - Clock speeds have not changed
  - The memory wall persists and multi-core places further pressure on this problem
- Microprocessor manufacturers are producing high-volume CPUs with 8-16 cores on-chip
  - New consumer chip Ryzen from AMD has up to 16 cores
  - SIMD/vector extensions – SSE (streaming SIMD extensions) and AVX (advanced vector extensions)
  - Also seeing multi-core in the embedded domain




# Why are Graphics Processors of interest?

- Graphics Processing Units
  - More than 65% of Americans played a video game in 2009
  - High-end - primarily used for 3-D rendering for videogame graphics and movie animation
  - Mid/low-end – primarily used for computer displays
  - Manufacturers include NVIDIA, AMD/ATI, Intel (embedded)
  - Very competitive commodities market




# Why GPUs? Performance: CPU vs GPU (FLOPS)



Source: Karl Rupp's website  
[www.karlrupp.net](http://www.karlrupp.net)



# Why GPUs? Performance/Power: CPU vs GPU




Source: Karl Rupp's website  
[www.karlrupp.net](http://www.karlrupp.net)

# A wide range of GPU applications

- 3D image analysis
- Adaptive radiation therapy
- Acoustics
- Astronomy
- Audio
- Automobile vision
- Bioinformatics
- Biological simulation
- Broadcast
- Cellular automata
- Fluid dynamics
- Computer vision
- Cryptography
- CT reconstruction
- Data mining
- Digital cinema / projections
- Electromagnetic simulation
- Equity trading
- Film
- Financial
- GIS
- Holographic cinema
- Intrusion detection
- Machine learning
- Mathematics research
- Military
- Mine planning
- Molecular dynamics
- MRI reconstruction
- Multispectral imaging
- N-body simulation
- Network processing
- Neural network
- Oceanographic research
- Optical inspection
- Particle physics
- Protein folding
- Quantum chemistry
- Ray tracing
- Radar
- Reservoir simulation
- Robotic vision / AI
- Robotic surgery
- Satellite data analysis
- Seismic imaging
- Surgery simulation
- Surveillance
- Ultrasound
- Video conferencing
- Telescope
- Video
- Visualization
- Wireless
- X-Ray

# GPU as a General Purpose Computing Platform

- Speedups are impressive and ever increasing!



**Genetic Algorithm**

**2600 X**



**Real Time Elimination  
of Undersampling Artifacts for Numerical Fluid Mechanics**

**2300 X**



**Lattice-Boltzmann Method  
for Numerical Fluid Mechanics**

**1840 X**




**Total Variation Modeling**

**1000 X**




**Fast Total Variation for  
Computer Vision**

**1000 X**



**Monte Carlo Simulation  
Of Photon Migration**

**1000 X**



**Stochastic Differential  
Equations**

**675 X**



**K-Nearest Neighbor  
Search**

**470 X**

Source: CUDA Zone at [www.nvidia.com/cuda/](http://www.nvidia.com/cuda/)


# Parallel Computing

- We want to achieve high performance (HPC)
  - Improve speed of computation
- Two important aspects we should consider for parallel computing:
  - Architecture
  - Parallel Programming

# Architecture of a GPU


- GPUs have a many-core architecture.
- Many-core:
  - architecture with hundreds or thousands of cores.

CPU: Cache heavy,  
focused on individual  
thread performance



Irregular data accesses  
More cache + Control

GPU: ALU heavy,  
massively parallel,  
throughput-oriented



Regular data accesses

# Parallel Programming Software

- Hardware vendors have tried to roll their own
  - NVIDIA's CUDA
  - AMD's CTM and Brook+
  - Intel Ct Technology
- Software vendors are developing new parallelization technology
  - Multi-core aware operating systems – Microsoft (BarrelFish), Apple (Grand Central Dispatch), others
  - Parallelizing compilers – Microsoft Visual, LLVM, Open64, Portland Group, IBM XLC, others
  - Portable frameworks for heterogeneous computing – Kronos/OpenCL

# So how do we start writing parallel programs?

- When we think of writing code, there are two ways of tackling this problem.
- When we have a shared memory model.
- When we have a distributed system model.

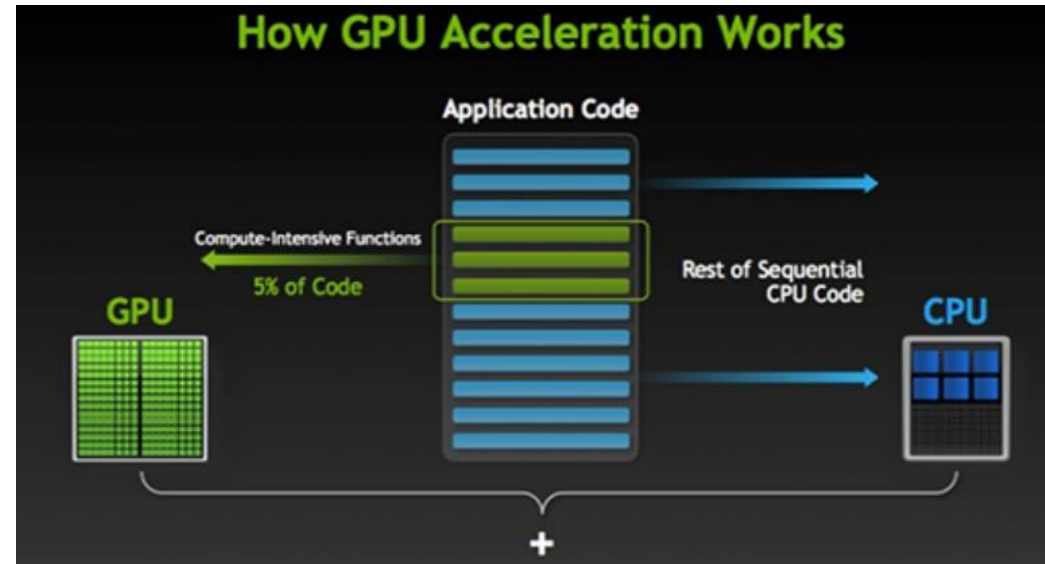
# So how do we start writing parallel programs?

- Shared memory system model
  - A single contiguous memory address space
  - Significantly reduces the burden of parallelizing programs
- Distributed systems model
  - Multiple memory address spaces are available
  - The programmer has the task of issuing explicit communication commands to manage synchronization of tasks and distributed memory




# Heterogeneous Architecture

- A heterogeneous application consists of two parts:
  - Host code. Runs on the CPU
  - Device code. Runs on the GPU



# Reasoning about Parallelism



# We need to start thinking in parallel

- To begin to utilize parallel systems such as multi-core CPUs, many-core GPUs and clusters, you need to understand parallelism
- We will explore this with a simple exercise
- We will explore some questions that will help you understand the challenges we must overcome to exploit parallel resources

# Jigsaw Puzzle

Henry Neeman, Lloyd Lee, Julia Mullen, and Gerard Newman. 2006. Analogies for teaching parallel computing to inexperienced programmers. *SIGCSE Bull.* 38, 4 (June 2006), 64-67. DOI=<http://dx.doi.org/10.1145/1189136.1189172>

# Jigsaw Puzzle

- Person A is working on a jigsaw puzzle
- It takes A 1 hour to complete the puzzle



# Jigsaw Puzzle

- Now, person B sits across the table from A and works with A on the puzzle.
- Lets assume that half the puzzle is grass and the other half is the sky.
- If A works on the grass and B on the sky, how long will it take the two of them to complete it?





# Jigsaw Puzzle

- If A and B both try to grab a piece out of the pile at the same time, will that slow them down?
- What if they grab for the same piece at the same time?



# Jigsaw Puzzle

- Can A and B work completely independently or will they need to work together at the horizon – that is, at the shared interface between their parts?





# Jigsaw Puzzle

- Now suppose that C and D sit at the table as well.
- Will there be more contention for the shared resource, or less, or the same?
- What about communication at the shared interfaces? How long will it take four people?



# Jigsaw Puzzle

- Now suppose that E, F, G and H sit at the table too.
- How long will it take the eight of them? If we keep adding people around the table, do we keep speeding up?



# Jigsaw Puzzle

- Parallelism isn't always the best way (at least, not straight forward).
- We need to understand the resources we have available and understand how to use them together in order to get a benefit in performance.



# Jigsaw Puzzle


- What if we have 8 puzzles to complete now?





# Parallelizing our lives

- Many of the tasks we perform in our everyday lives include significant parallelism
    - Can you name some of these?
  - How many of these activities could be carried out concurrently
    - Identify pairs of parallelizable activities
- Shower
  - Get Dressed
  - Eat breakfast
  - Wash our teeth
  - Dry your hair
  - Drive to work
  - Study
  - Check email
  - Use cellphone
  - Do exercise
  - Socialize
  - Etc...



# What is wrong with our world?




- Consider why many of these cannot presently be carried out in parallel
  - What would need to be changed in our physical world (e.g., showers, cars, Ipods) to allow us to complete many of these activities in parallel
  - How often is parallelism inhibited by our inability of carrying out two things at the same time?
- Estimate how much more quickly it would take to carry out these activities if you could change these physical systems

- Shower
- Get Dressed
- Eat breakfast
- Wash our teeth
- Dry your hair
- Drive to work
- Study
- Check email
- Use cellphone
- Do exercise
- Socialize
- Etc...



What is wrong with our world?  
Nothing!!



There is rampant parallelism in the natural world!

# Additional Resources

- This is a link to a very watchable and informative C++ playlist on Youtube:
  - <https://www.youtube.com/watch?v=tvC1WCdV1XU&list=PLAE85DE8440AA6B83>
- This website contains many C++ example problems and solutions:
  - [http://www.worldbestlearningcenter.com/index\\_files/cpp-tutorial-variables\\_datatypes\\_exercises.htm](http://www.worldbestlearningcenter.com/index_files/cpp-tutorial-variables_datatypes_exercises.htm)
- This is a link to a great linux playlist on youtube:
  - <https://www.youtube.com/watch?v=HjuHHI60s44&list=PL6gx4Cwl9DGCkg2uj3PxUWhMDuTw3VKjM>
- This website contains many Linux examples and tutorials:
  - <http://www.ee.surrey.ac.uk/Teaching/Unix/>
- And this website has an interactive terminal emulator that allows for safe practicing of commands:
  - <https://www.codecademy.com/learn/learn-the-command-line>



# Connecting to the Discovery Cluster

- How many of you have Windows?
  - Install FileZilla
  - Install Putty
- How many of you have Linux/Mac?
  - Use the Unix terminal
  - Use scp to copy files from/to the server.
    - Learn how to use this.