

Image Processing

Why Do We Process Images?

- Enhancement and restoration
 - Remove artifacts and scratches from an old photo/movie
 - Improve contrast and correct blurred images
- Transmission and storage
 - Images and Video can be more effectively transmitted and stored
- Information analysis and automated recognition
 - Recognizing terrorists
- Evidence
 - Careful image manipulation can reveal information not present
 - Detect image tampering
- Security and rights protection
 - Encryption and watermarking preventing illegal content manipulation

Some examples

Compression

- Color image of 600x800 pixels
 - Without compression
 - $600 \times 800 \times 24 \text{ bits/pixel}$
 $= 11\,520\text{K bits} = 1.44\text{M bytes}$
 - After JPEG compression (popularly used on web)
 - only 89K bytes
 - compression ratio $\sim 16:1$
- Movie
 - 720x480 per frame, 30 frames/sec, 24 bits/pixel
 - Raw video $\sim 243\text{M bits/sec}$
 - DVD $\sim \text{about } 5\text{M bits/sec}$
 - Compression ratio $\sim 48:1$



“Library of Congress” by M.Wu
(600x800)

Denoising



From X.Li

<http://www.ee.princeton.edu/~lixin/denoising.htm>

Deblurring



Blurred & noisy image




Restored image

From Mathworks


Visual Mosaicing

- Stitch photos together without thread or scotch tape





Visible Digital Watermarks

- from IBM Watson web page
“Vatican Digital Library”



Invisible Watermark

- 1st & 30th Mpeg4.5Mbps frame of original, marked, and their luminance difference
- human visual model for imperceptibility: protect smooth areas and sharp edges




Error Concealment

(a) original Lenna image



(b) corrupted Lenna image



(c) concealed Lenna image



25% blocks in a
checkerboard pattern are
corrupted

corrupted blocks are
concealed via edge-directed
interpolation

Examples were generated using the source codes provided by W.Zeng.


Image Super-Resolution

- • Super-Resolution(SR)
- Low resolution images \rightarrow High resolution images
- • Face Hallucination
 - Super-resolution on human faces
- • Constraints
 - Reconstruction constraints
 - Close to input image when blurred and down-sampled
 - Global structure constraints
 - Recovered image should be like a face
 - Sparsity assumption
 - Sparse coefficients are preserved during down-sampling



Courtesy of Jianchao Yang

Face Hallucination



Input
image

Bicubic
interpo-
lation

Back
projection

NMF

Sparse
Coding

Original

Courtesy of Jianchao Yang




What is an Image?

What is an image?

- We can think of an image as a function, f :
 - $f(x, y)$ gives the intensity at position (x, y)
 - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
 - $f: [a, b] \times [c, d] \rightarrow [0, 1]$
- A color image is just three functions pasted together. We can write this as a “vector-valued” function



$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

Images as functions



Digital Image

- In digital imaging, a pixel (picture element) is the smallest piece of information in an image.
- The word pixel is based on a contraction of pix (for "pictures") and el (for "element")
- Each pixel is a sample of an original image, where more samples typically provide a more accurate representation of the original.
- The intensity of each pixel is variable; in color systems, each pixel has typically three or four components such as red, green, and blue, or cyan, magenta, yellow, and black.



<http://en.wikipedia.org/wiki/Pixel>

Google.com image search

What is a digital image?

- We usually operate on digital (discrete) images:
 - Sample the 2D space on a regular grid
 - Quantize each sample (round to nearest integer)
- The image can now be represented as a matrix of integer values

$j \longrightarrow$

$i \downarrow$

62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

Other Applications

Image Enhancement

a b
c d

FIGURE 3.9
(a) Aerial image.
(b)–(d) Results of
applying the
transformation in
Eq. (3.2-3) with
 $c = 1$ and
 $\gamma = 3.0, 4.0,$ and
 5.0 , respectively.
(Original image
for this example
courtesy of
NASA.)


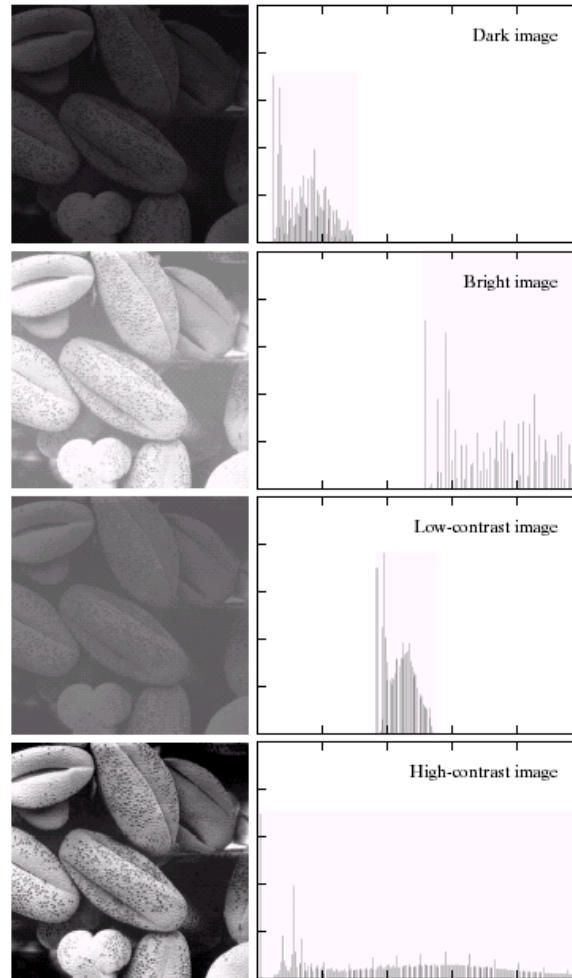



Image Histograms



a b

FIGURE 3.15 Four basic image types: dark, light, low contrast, high contrast, and their corresponding histograms. (Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Histogram Equalization



a b c

FIGURE 3.17 (a) Images from Fig. 3.15. (b) Results of histogram equalization. (c) Corresponding histograms.

How can we read
images?

Image file Formats

- Image file formats are standardized means of organizing and storing digital images.
- Types of files
 - JPEG
 - PNG
 - BMP
 - GIF
 - PPM
 - PGM
 - Etc



How can we read these formats

- So many different formats, which one is the best one?
- Which one is the easiest to interpret?
- The answer to these questions is not easy.
- Best way to read these types of files is through libraries such as Cimg, libjpeg, OpenCV.

Simplest version

- portable pixmap format (PPM), the portable graymap format (PGM)
 - These formats don't use any sort of compression so they are big files but easy to read the data.

```
P2
# Shows the word "FEEP" (example from Netpbm man page on PGM)
24 7
15
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 3 3 3 3 0 0 7 7 7 7 0 0 11 11 11 11 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 0 11 0 0 0 0 0 15 0 0 15 0
0 3 3 3 0 0 0 7 7 7 0 0 0 11 11 11 0 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 11 0 0 0 0 0 15 0 0 0 0
0 3 0 0 0 0 0 7 7 7 7 0 0 11 11 11 11 0 0 15 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```



Simplest version

- In this PGM example
 - P2 means PGM
 - The next two numbers give the width and the height.
 - The next number represents the maximum value (numbers of grey between black and white)
 - Black is 0 and max value is white.
 - Then follows the matrix with the pixel values.
 - The PGM and PPM formats (both ASCII and binary versions)

```
P2
# Shows the word "FEEP" (example from Netpbm man page on PGM)
24 7
15
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 3 3 3 3 0 0 7 7 7 7 0 0 11 11 11 11 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 0 11 0 0 0 0 0 15 0 0 15 0
0 3 3 3 0 0 0 7 7 7 0 0 0 0 11 11 11 0 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 0 11 0 0 0 0 0 15 0 0 0 0
0 3 0 0 0 0 0 7 7 7 7 0 0 0 11 11 11 11 0 0 15 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```



Simple algorithm

- Brightness
 - Increase/Decrease the intensity of the pixels by a specific value
 - Lets look at the code

What is the Makefile

```
# environment
SM := 35

# compilers
CUDA_PATH ?= $(shell test -d /shared/apps/cuda7.0 && echo /shared/apps/cuda7.0)
ifeq ($(CUDA_PATH),)
    CUDA_PATH := $(shell test -d /usr/local/cuda-7.0 && echo /usr/local/cuda-7.0)
endif

GCC := g++
NVCC := $(CUDA_PATH)/bin/nvcc
MPI = mpiCC

# libraries
CUDA_LIB_PATH := $(CUDA_PATH)/lib

# Remove function
RM = rm -f

# Compiler flags:
# -g debugging information
# -Wall turns on most compiler warnings
GENCODE_FLAGS := -gencode arch=compute_$(SM),code=sm_$(SM)
LIB_FLAGS := -lcudadevrt -lcudart
ifeq ($(OS), DARWIN)
    CCFLAGS := -stdlib=libstdc++
else
    CCFLAGS := -O3
endif
NVCCFLAGS :=
GccFLAGS = -fopenmp -O3
MPIFLAGS = -Wno-deprecated

debug: GccFLAGS += -DDEBUG -g -Wall
debug: MPIFLAGS += -DDEBUG
debug: NVCCFLAGS += -g -G
debug: all

# The build target executable:
TARGET = brightness

all: build

build: $(TARGET)

$(TARGET): lib/dlink.o lib/main.o lib/alg/$(TARGET).o lib/img/imghandler.o lib/alg/locationhandler.o
    $(NVCC) $(NVCCFLAGS) $^ -o $@ $(GENCODE_FLAGS) -link
lib/dlink.o: lib/alg/$(TARGET).o
    $(NVCC) $(NVCCFLAGS) $^ -o $@ $(GENCODE_FLAGS) -dlink
lib/main.o: lib/main.cpp lib/config.h
    $(GCC) $(GccFLAGS) -c $< -o $@
lib/alg/locationhandler.o: lib/alg/locationhandler.cpp lib/alg/locationhandler.h lib/config.h
    $(GCC) $(GccFLAGS) -c $< -o $@
lib/alg/$(TARGET).o: lib/alg/$(TARGET).cu lib/config.h
    $(NVCC) $(NVCCFLAGS) -dc $< -o $@ $(GENCODE_FLAGS)
lib/img/imghandler.o: lib/img/imghandler.cpp lib/config.h
    $(GCC) $(GccFLAGS) -c $< -o $@

clean:
    $(RM) $(TARGET) *.o lib/*.o *.tar* *.core*
```


What is the Makefile

- Makefiles are a simple way to organize code compilation.
- A makefile is a file containing a set of directives used with the make build automation tool from GNU.
- Incredibly useful when your project becomes bigger and bigger.

```
# environment
SM := 35

# compilers
CUDA_PATH ?= $(shell test -d /shared/apps/cuda7.0 && echo /shared/apps/cuda7.0)
ifeq ($(CUDA_PATH),)
  CUDA_PATH := $(shell test -d /usr/local/cuda-7.0 && echo /usr/local/cuda-7.0)
endif

GCC := g++
NVCC := $(CUDA_PATH)/bin/nvcc
MPI = mpiCC

# libraries
CUDA_LIB_PATH := $(CUDA_PATH)/lib

# Remove function
RM = rm -f

# Compiler flags:
# -g debugging information
# -Wall turns on most compiler warnings
GENCODE_FLAGS := -gencode arch=compute_$(SM),code=sm_$(SM)
LIB_FLAGS := -lcudadevrt -lcudart
ifeq ($(OS), DARWIN)
  CCFLAGS := -stdlib=libstdc++
else
  CCFLAGS := -O3
endif
NVCCFLAGS :=
GccFLAGS = -fopenmp -O3
MPIFLAGS = -Wno-deprecated

debug: GccFLAGS += -DDEBUG -g -Wall
debug: MPIFLAGS += -DDEBUG
debug: NVCCFLAGS += -g -G
debug: all

# The build target executable:
TARGET = brightness

all: build

build: $(TARGET)

$(TARGET): lib/dlink.o lib/main.o lib/alg/$(TARGET).o lib/img/imghandler.o lib/alg/locationhandler.o
  $(NVCC) $(NVCCFLAGS) $^ -o $@ $(GENCODE_FLAGS) -link

lib/dlink.o: lib/alg/$(TARGET).o
  $(NVCC) $(NVCCFLAGS) $^ -o $@ $(GENCODE_FLAGS) -dlink

lib/main.o: lib/main.cpp lib/config.h
  $(GCC) $(GccFLAGS) -c $< -o $@

lib/alg/locationhandler.o: lib/alg/locationhandler.cpp lib/alg/locationhandler.h lib/config.h
  $(GCC) $(GccFLAGS) -c $< -o $@

lib/alg/$(TARGET).o: lib/alg/$(TARGET).cu lib/config.h
  $(NVCC) $(NVCCFLAGS) -dc $< -o $@ $(GENCODE_FLAGS)

lib/img/imghandler.o: lib/img/imghandler.cpp lib/config.h
  $(GCC) $(GccFLAGS) -c $< -o $@

clean:
  $(RM) $(TARGET) *.o lib/*.o *.tar* *.core*
```

```

int main(int argc, char* argv[]) {

    // Files needed
    char* imageFile = NULL;

    // Load Intensity Image
    image<unsigned char*> intensityInput = loadPGM(imageFile);

    cpu.height = intensityInput->height();
    cpu.width  = intensityInput->width();

    cpu.gridXSize = 1 + ((cpu.width - 1) / TILE_SIZE);
    cpu.gridYSize = 1 + ((cpu.height - 1) / TILE_SIZE);

    int XSize = cpu.gridXSize*TILE_SIZE;
    int YSize = cpu.gridYSize*TILE_SIZE;

    cpu.size = XSize*YSize;

    cpu.intensity = new unsigned char[cpu.size];

    for (unsigned int y = 0 ; y < YSize ; y++){
        for (unsigned int x = 0 ; x < XSize ; x++){
            unsigned int newLocation =
imageLocation (x, y, cpu.gridXSize);
            if (x < cpu.width && y < cpu.height) {
                cpu.intensity[newLocation] =
intensityInput->data[y*cpu.width + x];
            } else{
                // Necessary in case image size is
not a multiple of TILE_SIZE
                cpu.intensity[newLocation] = 0;
            }
        }
    }

    cpu.result = brightness(cpu.intensity,
                            cpu.height,
                            cpu.width);

```

```

// Output RGB images
char filename[64];
sprintf(filename, "result.ppm");

srand(1000);

Color color;

// Create output image
image<Color> output = image<Color>(cpu.width, cpu.height,
true);
image<Color*> im = &output;

Color randomcolor = randomColor();
for (unsigned int y = 0 ; y < cpu.height ; y++){
    for (unsigned int x = 0 ; x < cpu.width ; x++){
        unsigned int newLocation = imageLocation
(x, y, cpu.gridXSize);

        color.r = cpu.result[newLocation];
        color.g = cpu.result[newLocation];
        color.b = cpu.result[newLocation];
        im->access[y][x] = color;
    }
}
savePPM(im, filename);

// Free resources and end the program
free(cpu.intensity);

return 0;
}

```

```

unsigned char *brightness(unsigned char *intensity,
                          unsigned int height,
                          unsigned int width){

    int gridSize = 1 + ((width - 1) / TILE_SIZE);
    int gridYSize = 1 + ((height - 1) / TILE_SIZE);

    int XSize = gridSize*TILE_SIZE;
    int YSize = gridYSize*TILE_SIZE;

    // Both are the same size (CPU/GPU).
    gpu.size = XSize*YSize;

    // Allocate arrays in GPU memory

    checkCuda(cudaMalloc((void**)&gpu.intensity
, gpu.size*sizeof(char)));
    checkCuda(cudaMalloc((void**)&gpu.result
, gpu.size*sizeof(char)));

    // Allocate result array in CPU memory
    gpu.resultOnCPU = new unsigned char[gpu.size];

    checkCuda(cudaMemcpy(gpu.intensity,
                          intensity,
                          gpu.size*sizeof(char),
                          cudaMemcpyHostToDevice));

    checkCuda(cudaDeviceSynchronize());

    dim3 dimGrid(gridXSize, gridYSize);
    dim3 dimBlock(BLOCK_TILE_SIZE, BLOCK_TILE_SIZE);

    // Launch kernel to begin image segmenation
    brightnessAlgorithm<<<dimGrid,
dimBlock>>>(gpu.intensity,
                                                    gpu.result,

```

```

                                                    inc);

    checkCuda(cudaDeviceSynchronize());

    // Retrieve results from the GPU
    checkCuda(cudaMemcpy(gpu.resultOnCPU,
                          gpu.result,
                          gpu.size*sizeof(char),
                          cudaMemcpyDeviceToHost));

    // Free resources and end the program
    checkCuda(cudaFree(gpu.intensity));
    checkCuda(cudaFree(gpu.result));

    return(gpu.resultOnCPU);

```

```
__global__ void brightnessAlgorithm(unsigned char
*intensity,
                                   unsigned char
*result,
                                   unsigned int inc){

    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int bx = blockIdx.x;
    int by = blockIdx.y;

    // Read Input Data
    int x = bx*TILE_SIZE+tx;
    int y = by*TILE_SIZE+ty;

    int location = y*(gridDim.x*TILE_SIZE)+x;

    unsigned char value = intensity[location];

    // Algorithm
    if (value + inc > 255) result[location] =
255;
    else result[location] = value + inc;

}
```


How can we improve brightness?


- Shared memory?
- Constant Memory?
- Texture memory?

How can we improve brightness?

- More processing per thread
 - Adding more work to each thread
- Bring in more data on each read

A more complex example

- Sobel algorithm
 - Sobel edge detection
 - Find the boundaries of the image where there is significant difference as compared to neighboring “pixels” and replace values to find edges.



Sobel Algorithm

- Looks for edges in both horizontal and vertical directions, then combine the information into a single metric.
- The masks are:

$$y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$


$$x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Edge Magnitude =

$$\sqrt{x^2 + y^2}$$

Edge Direction =

$$\tan^{-1} \left[\frac{y}{x} \right]$$



Sobel Algorithm

Input



Output



```

• int main(int argc, char* argv[]) {
•
•     // Files needed
•     char* imageFile = NULL;
•
•     // Load Intensity Image
•     image<unsigned char*> intensityInput =
loadPGM(imageFile);
•
•     cpu.height = intensityInput->height();
•     cpu.width  = intensityInput->width();
•
•     cpu.gridXSize = 1 + ((cpu.width - 1) / TILE_SIZE);
•     cpu.gridYSize = 1 + ((cpu.height - 1) / TILE_SIZE);
•
•     int XSize = cpu.gridXSize*TILE_SIZE;
•     int YSize = cpu.gridYSize*TILE_SIZE;
•
•     cpu.size = XSize*YSize;
•
•     cpu.intensity = new unsigned char[cpu.size];
•
•     for (unsigned int y = 0 ; y < YSize ; y++){
•         for (unsigned int x = 0 ; x < XSize ; x++){
•             unsigned int newLocation =
imageLocation (x, y, cpu.gridXSize);
•             if (x < cpu.width && y < cpu.height) {
•                 cpu.intensity[newLocation] =
intensityInput->data[y*cpu.width + x];
•             } else{
•                 // Necessary in case image size
is not a multiple of TILE_SIZE
•                 cpu.intensity[newLocation] = 0;
•             }
•         }
•     }
•
•     cpu.result = sobel(cpu.intensity,
•                         cpu.height,

```

```

•                         cpu.width);
•
•     // Output RGB images
•     char filename[64];
•     sprintf(filename, "result.ppm");
•
•     srand(1000);
•
•     Color color;
•
•     // Create output image
•     image<Color> output = image<Color>(cpu.width, cpu.height,
true);
•     image<Color*> im = &output;
•
•     Color randomcolor = randomColor();
•     for (unsigned int y = 0 ; y < cpu.height ; y++){
•         for (unsigned int x = 0 ; x < cpu.width ; x++){
•             unsigned int newLocation = imageLocation
(x, y, cpu.gridXSize);
•
•             color.r = cpu.result[newLocation];
•             color.g = cpu.result[newLocation];
•             color.b = cpu.result[newLocation];
•             im->access[y][x] = color;
•         }
•     }
•     savePPM(im, filename);
•
•     // Free resources and end the program
•     free(cpu.intensity);
•
•     return 0;
• }

```

```

unsigned char *sobel(unsigned char *intensity,
                    unsigned int height,
                    unsigned int width){

    int gridSize = 1 + (( width - 1) /
TILE_SIZE);
    int gridSize = 1 + ((height - 1) /
TILE_SIZE);

    int XSize = gridSize*TILE_SIZE;
    int YSize = gridSize*TILE_SIZE;

    // Both are the same size (CPU/GPU).
    gpu.size = XSize*YSize;

    // Allocate arrays in GPU memory
    checkCuda(cudaMalloc((void**)&gpu.intensity
, gpu.size*sizeof(char)));
    checkCuda(cudaMalloc((void**)&gpu.result
, gpu.size*sizeof(char)));

    // Allocate result array in CPU memory
    gpu.resultOnCPU = new unsigned
char[gpu.size];

    checkCuda(cudaMemcpy(gpu.intensity,
                        intensity,
                        gpu.size*sizeof(char),
                        cudaMemcpyHostToDevice));

    checkCuda(cudaDeviceSynchronize());

```

```

        dim3 dimGrid(gridXSize, gridYSize);
        dim3 dimBlock(BLOCK_TILE_SIZE,
BLOCK_TILE_SIZE);

        // Launch kernel to begin image segmenation
        sobelAlgorithm<<<dimGrid,
dimBlock>>>(gpu.intensity,
gpu.result,threshold);

        checkCuda(cudaDeviceSynchronize());

        // Retrieve results from the GPU
        checkCuda(cudaMemcpy(gpu.resultOnCPU,
                        gpu.result,
                        gpu.size*sizeof(char),
                        cudaMemcpyDeviceToHost));

        // Free resources and end the program
        checkCuda(cudaFree(gpu.intensity));
        checkCuda(cudaFree(gpu.result));

        return(gpu.resultOnCPU);

}

```

```

__global__ void sobelAlgorithm(unsigned char
*intensity,
                                unsigned char
*result,
                                unsigned int
threshold) {

    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int bx = blockIdx.x;
    int by = blockIdx.y;

    int x = bx*TILE_SIZE+tx;
    int y = by*TILE_SIZE+ty;

    int xsize = TILE_SIZE*gridDim.x;
    int ysize = TILE_SIZE*gridDim.y;

    if (y > 1 && y < ysize-1 && x > 1 && x
< xsize-1) {
        int location = y*xsize+x;
        int sum1 = input[ xsize * (y-1) + x+1 ] -
                    input[ xsize * (y-1) + x-1 ] +
                    2 * input[ xsize * (y)    + x+1 ] -
                    2 * input[ xsize * (y)    + x-1 ] +

```

```

                    input[ xsize * (y+1) + x+1 ] -
                    input[ xsize * (y+1) + x-1 ];
        int sum2 = input[ xsize * (y-1) + x-1 ] +
                    2 * input[ xsize * (y-1) + x    ] +
                    input[ xsize * (y-1) + x+1 ] -
                    input[ xsize * (y+1) + x-1 ] -
                    2 * input[ xsize * (y+1) + x    ] -
                    input[ xsize * (y+1) + x+1 ];

        int magnitude =
sum1*sum1+sum2*sum2;

        if (magnitude > threshold)
            result[location] =
255;

        else
            result[location] = 0;

    }
}

```