

GPU Programming

(in Cuda)

Julian Gutierrez

NUCAR

Session 11

Outline

- Applications
 - Parallel Reduction
 - Prefix Sum (Scan)
 - Histogram
 - Convolution

Parallel Reduction

- A popular class of computation
- Goal: To master the concept of control divergence through reduction trees (how to be work efficient).

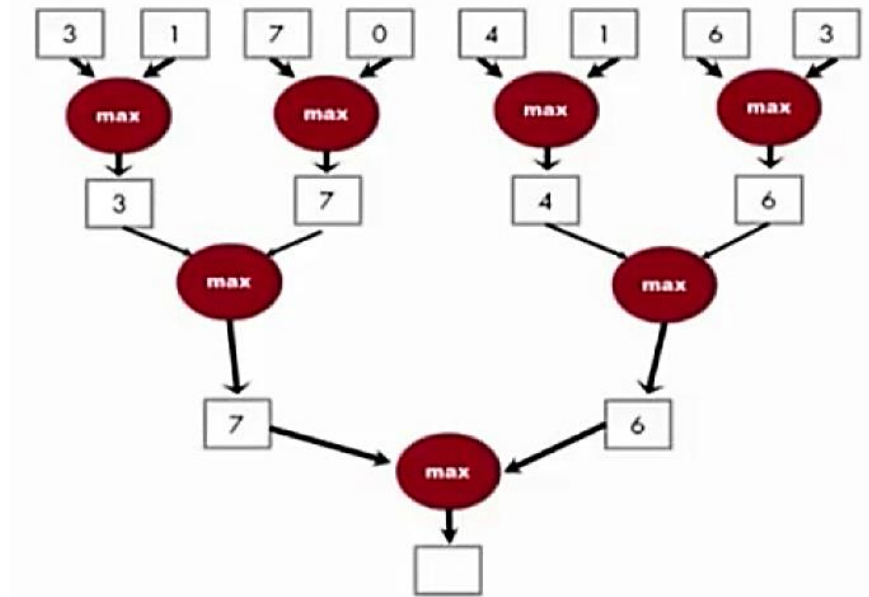
Parallel Reduction

- Example
- Calculate Max Value from vector:



Parallel Reduction

- Example
- Calculate Max Value from vector:
- A parallel reduction tree algorithm performs $N-1$ operations in $\log(N)$ steps



Parallel Reduction

- A commonly used strategy for processing large input data sets
 - There is no required order of processing elements in a data set (associative and commutative)
 - Partition the data set into smaller chunks
 - Have each thread to process a chunk
 - Use a reduction tree to summarize the results from each chunk into the final answer
- Google and Hadoop MapReduce frameworks are examples of this pattern
- We will focus on the reduction tree step for now.

Parallel Reduction

- Example: Handling privatization
 - Multiple threads write into an output location
 - Replicate the output location so that each thread has a private output location
 - Use a reduction tree to combine the values of private locations into the original output location

Parallel Reduction

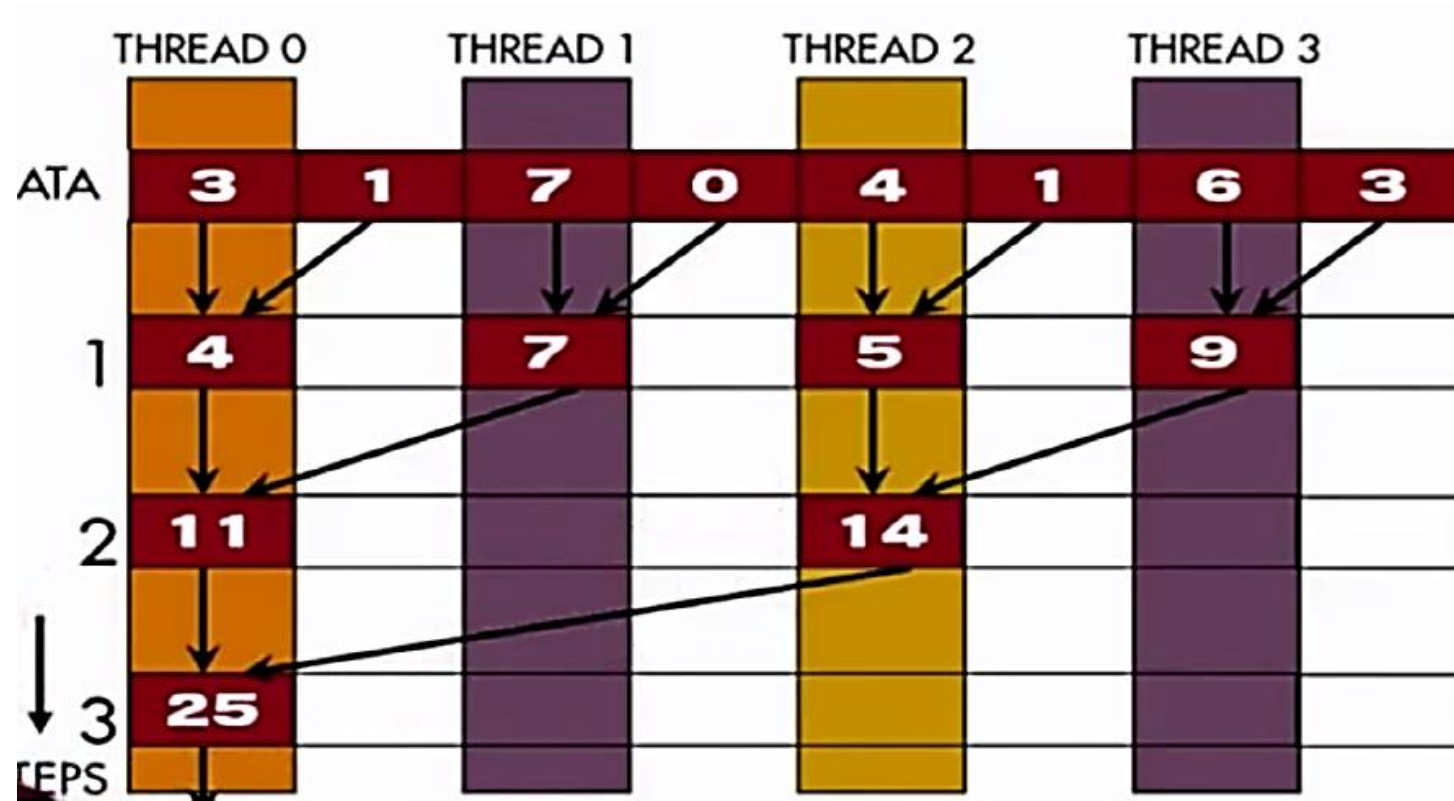
- Use: Summarize a set of input values into one value using a “reduction operation”
 - Max
 - Min
 - Sum
 - Product
- You need to initialize the result as an identity value for the reduction operation
 - Smallest possible value for max reduction
 - Largest possible value for min reduction
 - 0 for sum reduction
 - 1 for product reduction

Parallel Reduction

- For N input values, the reduction tree performs:
 - $(1/2)N + (1/4)N + (1/8)N + \dots (1/N) = N-1$ operations
 - In $\log(N)$ steps – 1 000 000 input values take 20 steps
 - Assuming that we have enough execution resources
 - Average parallelism $(N-1)/\log(N)$
 - For $N = 1\,000\,000$ average parallelism is 50 000
 - However, peak resource requirement is 500 000!
- This is a work-efficient parallel algorithm
 - The amount of work done is comparable to sequential
 - Many parallel algorithms are not work efficient

Parallel Reduction

- An example



Parallel Reduction

- An example
 - Each thread block takes $2 \times \text{BlockDim}$ input elements
 - Each threads loads 2 elements into shared memory

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
Unsigned int start = 2*blockIdx.x*blockDim.x;
partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start+ blockDim.x+t];
```

Parallel Reduction


```
for (unsigned int stride = 1;
     stride <= blockDim.x; stride *= 2)
{
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

- Why do we need `__syncthreads()`?

- THREAD 0 IN EACH THREAD BLOCK WRITE THE SUM OF THE THREAD BLOCK IN `PARTIALSUM[0]` INTO A VECTOR INDEXED BY THE `BLOCKIDX.X`
- THERE CAN BE A LARGE NUMBER OF SUCH SUMS IF THE ORIGINAL VECTOR IS VERY LARGE
 - THE HOST CODE MAY ITERATE AND LAUNCH ANOTHER KERNEL
- IF THERE ARE ONLY A SMALL NUMBER OF SUMS, THE HOST CAN SIMPLY TRANSFER THE DATA BACK AND ADD THEM TOGETHER

Parallel Reduction

Parallel Reduction: Interleaved Addressing




New Problem: Shared Memory Bank Conflicts

Parallel Reduction

- A better reduction?

Parallel Reduction: Sequential Addressing



Sequential addressing is conflict free

Parallel Reduction

- There's a problem

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

Parallel Reduction

- There's a problem
 - Idle threads!
 - Half of the threads are idle on first loop iteration!
 - This is wasteful....

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```


Parallel Reduction

- Halve the number of blocks, and replace single load:

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
sdata[tid] = g_idata[i];
__syncthreads();
```

- With two loads and first add of the reduction:

```
// perform first level of reduction,
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

Parallel Reduction

- Unrolling the last warp
 - As reduction proceeds, # “active” threads decreases
 - When $s \leq 32$, we have only one warp left
 - Instructions are SIMD synchronous within a warp
 - That means when $s \leq 32$:
 - We don't need to `__syncthreads()`
 - We don't need “if (tid < s)” because it doesn't save any work
 - Lets unroll the last 6 iterations of the inner loop

Parallel Reduction

- Note: This saves useless work in all warps, not just the last one!
- Without unrolling, all warps execute every iteration of the for loop and if statement

```
__device__ void warpReduce(volatile int* sdata, int tid) {  
    sdata[tid] += sdata[tid + 32];  
    sdata[tid] += sdata[tid + 16];  
    sdata[tid] += sdata[tid + 8];  
    sdata[tid] += sdata[tid + 4];  
    sdata[tid] += sdata[tid + 2];  
    sdata[tid] += sdata[tid + 1];  
}
```

↑

IMPORTANT:
For this to be correct,
we must use the
“volatile” keyword!

```
// later...  
for (unsigned int s=blockDim.x/2; s>32; s>>=1) {  
    if (tid < s)  
        sdata[tid] += sdata[tid + s];  
    __syncthreads();  
}  
  
if (tid < 32) warpReduce(sdata, tid);
```

Parallel Reduction

```
template <unsigned int blockSize>
__device__ void warpReduce(volatile int *sdata, unsigned int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}

template <unsigned int blockSize>
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n) {
    extern __shared__ int sdata[];
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*(blockSize*2) + tid;
    unsigned int gridSize = blockSize*2*gridDim.x;
    sdata[tid] = 0;

    while (i < n) { sdata[tid] += g_idata[i] + g_idata[i+blockSize]; i += gridSize; }
    __syncthreads();

    if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads(); }
    if (blockSize >= 256) { if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads(); }
    if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads(); }

    if (tid < 32) warpReduce(sdata, tid);
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```



Final Optimized Kernel

Prefix Sum (Scan)

- Definition

- Prefix sum, cumulative sum, inclusive scan, or simply scan of a sequence of numbers x_0, x_1, x_2, \dots . Gives a second sequence of numbers y_0, y_1, y_2, \dots , the sums of prefixes (running totals) of the input sequence.
- Any binary operation (not just the addition operation).

- Example

- If (+) is addition, then scan on the set: `[3 1 7 0 4 1 6 3]`
 - Returns the set: `[0 3 4 11 11 15 16 22]`

Note: Exclusive scan: last input element is not included in the result

Prefix Sum (Scan)

- A Naïve inclusive parallel Scan
 - Assign one thread to calculate each y element
 - Have every thread add up all x elements needed for the y element

$$y0 = x0$$

$$y1 = x0 + x1$$

$$y2 = x0 + x1 + x2$$

Prefix Sum (Scan)


- A Naïve inclusive parallel Scan
 - Assign one thread to calculate each y element
 - Have every thread add up all x elements needed for the y element

$$\begin{aligned}y0 &= x0 \\y1 &= x0 + x1 \\y2 &= x0 + x1 + x2\end{aligned}$$

- Note: Parallel programming is easy as long as you do not care about performance


Prefix Sum (Scan)

- Compared to naïve version, a slight better version – Hillis Steele Scan.
- Iterate $\log(n)$ times: Threads stride to n : Add pairs of elements stride elements apart. Double stride at each iteration. (note must double buffer shared memory arrays)




Prefix Sum (Scan)

- Compared to naïve version, a slight better version – Hillis Steele Scan.
- Iterate $\log(n)$ times: Threads stride to n : Add pairs of elements stride elements apart. Double stride at each iteration. (note must double buffer shared memory arrays)
- This scan algorithm is not that work efficient
 - Sequential scan algorithm does $n-1$ adds
 - How many does this one do?
 - What happens if the # of elements is 10^6 ?



Prefix Sum (Scan)

- Even better one: Blelloch Scan.
- Reduces in $\log(n)$, downsweeps in $\log(n)$.
- Total steps: $2 \log(n)$
- This scan algorithm is work efficient




Prefix Sum (Scan)

- Useful in implementation of several parallel algorithms:
 - Radix sort
 - Quicksort
 - String comparison
 - Lexical analysis
 - Stream compaction
 - Polynomial evaluation
 - Solving recurrences
 - Tree operations
 - histograms
- Examples
 - Assigning camp slots
 - Assigning farmer market space
 - Allocating memory to parallel threads
 - Allocating memory buffer for communication channels

Histogram

- A histogram is a graphical representation of the distribution of numerical data.
- Bar Graph




Histogram

- A histogram is a graphical representation of the distribution of numerical data.
- Bar Graph
- Serial Algorithm

```
For (i=0; I < BIN_COUNT; i++)  
    result[i]=0;
```

```
For (i=0; I < measurements.size(); i++)  
    result[computeBin(measurements[i])]++;
```



Histogram

- Naïve implementation:

```
__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}
```

Histogram


- Naïve implementation:

```
__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}
```

WHY THE OBVIOUS METHOD DOESN'T WORK

THREAD:

- (1) READ BIN VALUE
- (2) INCREMENT BIN VALUE
- (3) WRITE BIN VALUE



Histogram

- Naïve implementation using atomic operations (Method 1):
 - This will avoid RAW hazards.

```
__global__ void simple_histo(int *d_bins, const int *d_in)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    atomicAdd(&(d_bins[myBin]), 1);
}
```


Histogram

- Naïve implementation using atomic operations (Method 1):
 - This will avoid RAW hazards.

```
__global__ void simple_histo(int *d_bins, const int *d_in)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    atomicAdd(&(d_bins[myBin]), 1);
}
```

QUIZ

- Histogram with 1M elements
- you can choose # of bins:



Histogram

- Redefining the method. Local histogram + reduction (Method 2)

PER-THREAD PRIVATIZED (LOCAL) HISTOGRAMS, THEN REDUCE
128 ITEMS · 8 THREADS · 3 BINS
(EACH THREAD GETS 16 ITEMS)



Histogram

- Redefining the method. Local histogram + reduction (Method 2)

PER-THREAD PRIVATIZED (LOCAL) HISTOGRAMS, THEN REDUCE
128 ITEMS · 8 THREADS · 3 BINS
(EACH THREAD GETS 16 ITEMS)



Q: DO WE NEED ATOMICS TO MANAGE ACCESS TO THESE LOCAL PER-THREAD HISTOGRAMS?

YES ☐
NO ☐

Histogram

- Redefining the method. Sort then reduce by key (Method 3)

Sort, then reduce by key


8 ENTRIES

3 DWS

(0, 1, 2)



(1) SORT



Histogram


- Final thoughts on histogram:
 - Using Atomic operations
 - Using per-thread histograms, and then reduce
 - Sort, then reduce by key
- Question
 - 256 threads, 8 bins. How many atomic adds are needed?
 - Atomic Technique.
 - Processing 16 elements per thread with local histogram, and then atomics.

Convolution

- Applications

- A popular array operation that is used in various forms in signal processing, digital recording, image processing, video processing, and computer vision.
- Convolution is often performed as a filter that transforms signals and pixels into more desirable values
 - Some filters smooth out the signal values so that one can see the big-picture trend
 - Others like Gaussian filters can be used to sharpen boundaries and edges of objects in images

Convolution




Convolution

- An array operation where each output data element is a weighted sum of a collection of neighboring input elements
- The weights used in the weighted sum calculation are defined by an input mask array, commonly referred to as the convolution kernel
 - We will refer to these mask arrays as convolution masks to avoid confusion.
 - The same convolution mask is typically used for all elements of the array.

Convolution

- 1D Convolution example
 - Commonly used for audio processing
 - Mask size is usually an odd number of elements for symmetry



Convolution

- Definition

- For a causal discrete-time FIR filter of order N, each value of the output sequence is a weighted sum of the most recent input values:


$$\begin{aligned}y[n] &= b_0x[n] + b_1x[n-1] + \cdots + b_Nx[n-N] \\ &= \sum_{i=0}^N b_i \cdot x[n-i],\end{aligned}$$

- Where:

- X[n] is the input signal
- Y[n] is the output signal,
- N is the filter order, an Nth-order filter has (N+1) terms on the right-hand side
- Bi is the value of the impulse response at the ith instance for $0 \leq i \leq N$ of an Nth-order filter. If the filter is a direct form FIR filter then bi is also a coefficient of the filter.

Convolution

- Calculation of output elements near the boundaries (beginning and end) of the input array need to deal with “ghost” elements
 - Different policies (0, replicates of boundary values, etc.)



Convolution

- This kernel forces all elements outside the vector range to 0


```
__global__ void convolution_1D_basic_kernel(float *N, float *M, float *P,
int Mask_Width, int Width) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;


    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);
    for (int j = 0; j < Mask_Width; j++) {
        if (N_start_point + j >= 0 && N_start_point + j < Width) {
            Pvalue += N[N_start_point + j]*M[j];
        }
    }
    P[i] = Pvalue;
}
```

Convolution

- 2D convolution




- 2D convolution – ghost cells




Convolution

- Tiled Convolution for 1D



Convolution


- Loading the left halo



```
int n = Mask_Width/2;
int halo_index_left = (blockIdx.x - 1)*blockDim.x + threadIdx.x;
if (threadIdx.x >= blockDim.x - n) {
    N_ds[threadIdx.x - (blockDim.x - n)] =
        (halo_index_left < 0) ? 0 : N[halo_index_left];
}
```

Convolution


- Loading the internals



```
N_ds[n + threadIdx.x] = N[blockIdx.x*blockDim.x + threadIdx.x];
```


Convolution

- Loading the right halo



```
int halo_index_right = (blockIdx.x + 1)*blockDim.x + threadIdx.x;
if (threadIdx.x < n) {
    N_ds[n + blockDim.x + threadIdx.x] =
        (halo_index_right >= Width) ? 0 : N[halo_index_right];
}
```

Convolution

```
__global__ void convolution_1D_basic_kernel(float *N, const float __restrict__ *M,
                                           float *P, int Mask_Width, int Width) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    __shared__ float N_ds[TILE_SIZE + MAX_MASK_WIDTH - 1];

    int n = Mask_Width/2;

    int halo_index_left = (blockIdx.x - 1)*blockDim.x + threadIdx.x;
    if (threadIdx.x >= blockDim.x - n) {
        N_ds[threadIdx.x - (blockDim.x - n)] =
            (halo_index_left < 0) ? 0 : N[halo_index_left];
    }

    N_ds[n + threadIdx.x] = N[blockIdx.x*blockDim.x + threadIdx.x];


    int halo_index_right = (blockIdx.x + 1)*blockDim.x + threadIdx.x;
    if (threadIdx.x < n) {
        N_ds[n + blockDim.x + threadIdx.x] =
            (halo_index_right >= Width) ? 0 : N[halo_index_right];
    }

    __syncthreads();

    float Pvalue = 0;
    for(int j = 0; j < Mask_Width; j++) {
        Pvalue += N_ds[threadIdx.x + j]*M[j];
    }
    P[i] = Pvalue;
}
```

Convolution


- Shared memory data reuse



Do we want to use Shared memory?

Convolution

- Shared memory data reuse



Element 2 is used by thread 4 (1X)

Element 3 is used by threads 4, 5 (2X)

Element 4 is used by threads 4, 5, 6 (3X)

Element 5 is used by threads 4, 5, 6, 7 (4X)

Element 6 is used by threads 4, 5, 6, 7 (4X)

Element 7 is used by threads 5, 6, 7 (3X)

Element 8 is used by threads 6, 7 (2X)

Element 9 is used by thread 7 (1X)

Do we want to use Shared memory?